Language Understanding

14 - Information Extraction and Knowledge Graphs

Hossein Zeinali

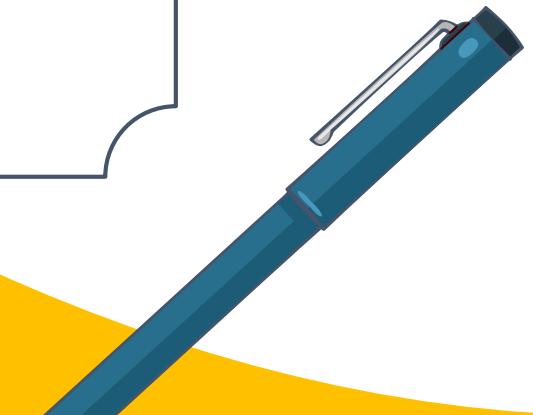

What is Knowledge Graph?

- *"A huge knowledge graph of interconnected entities and their attributes".* Amit Singhal, Senior Vice President at Google
- *"A knowledge based used by Google to enhance its search engine's results with semantic-search information gathered from a wide variety of sources"* http://en.wikipedia.org/wiki/Knowledge_Graph
- Based on information derived from many sources including *Freebase, CIA World Factbook, Wikipedia*
- By May 2020, contains 5 billion entities and more than 500 billion facts about and relationships between these different objects

Who Constructs Knowledge Graphs?

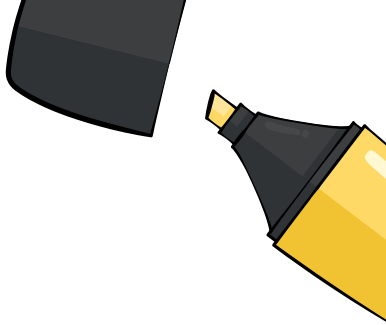

- Google
- Amazon
- Microsoft
- Bloomberg (business intelligence)
- National Cancer Institute (part of NIH)
- National Center for Biomedical Ontology (funded by NIH)
- Specialized knowledge graphs:
  - Health / life sciences (health-lifesci.schema.org)
  - Earth science
  - Agriculture
  - NCI Enterprise Vocabulary Services (evs.nci.nih.gov)

Google Knowledge Graph (GKG)

GKG enhances Google Search in three main ways:

- **Find the right thing:** Deals with the ambiguity of the language

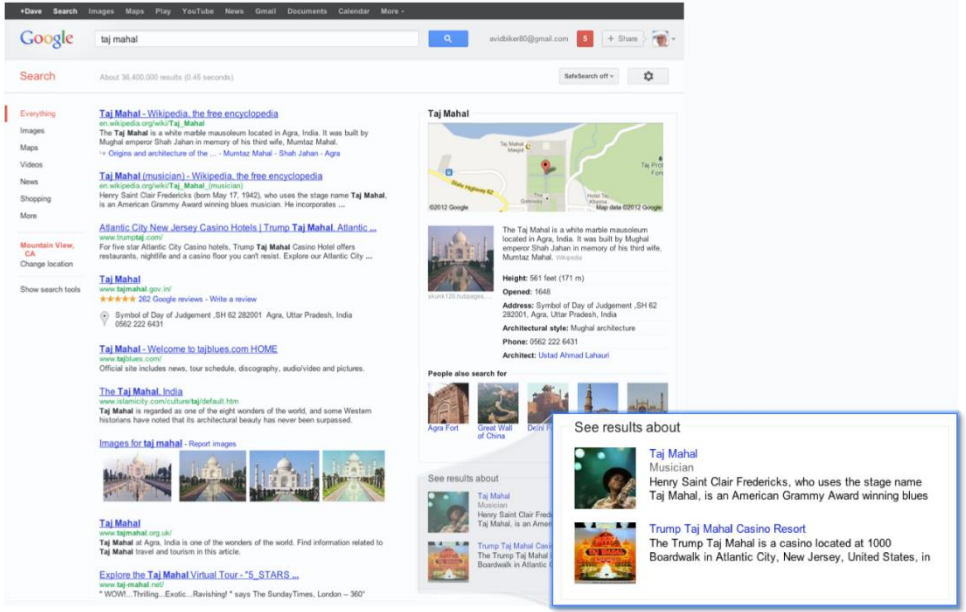

Google Knowledge Graph (GKG)

GKG enhances Google Search in three main ways:

- **Summaries:** summarize relevant content around that topic, including key facts about the entity

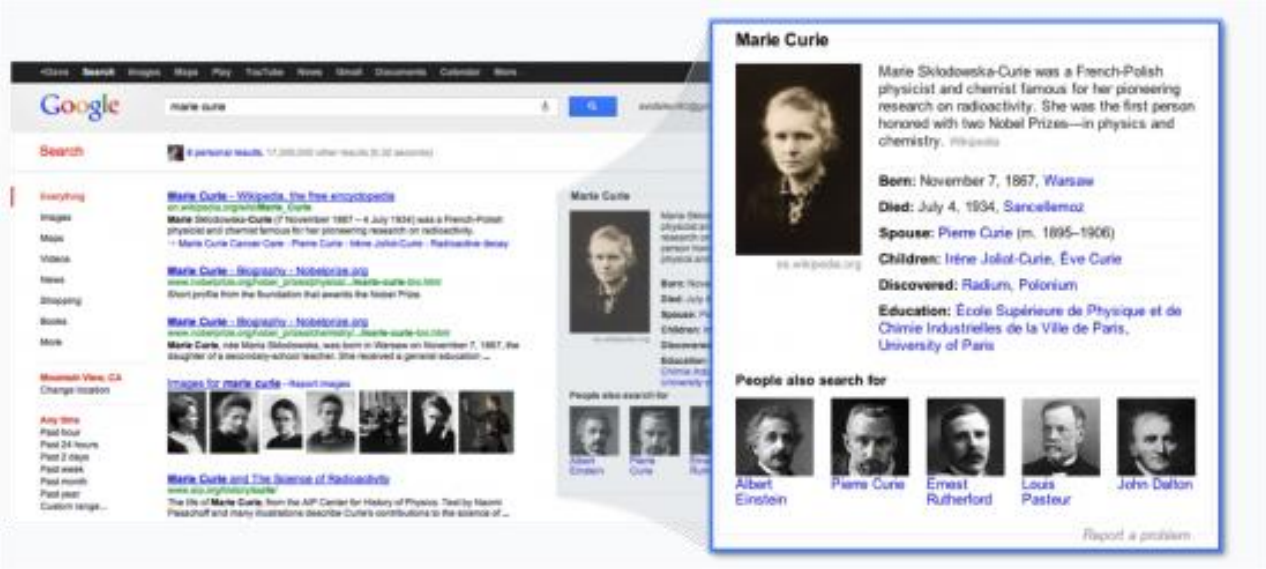

Google Knowledge Graph (GKG)

GKG enhances Google Search in three main ways:

- **Deeper and broader information:** reveal new facts, anticipate what the next questions and provide the information beforehand

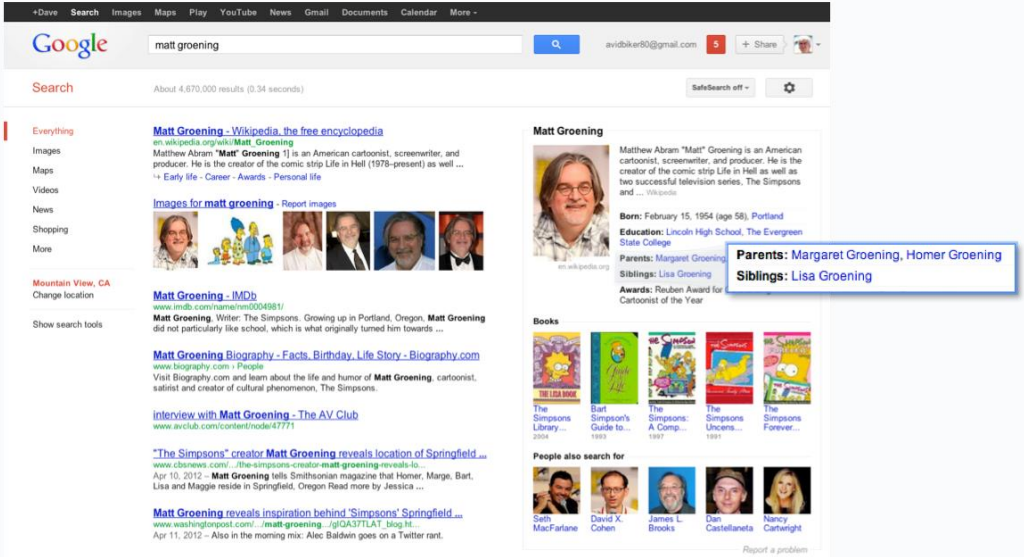

How GKG is used?

- Explore your search

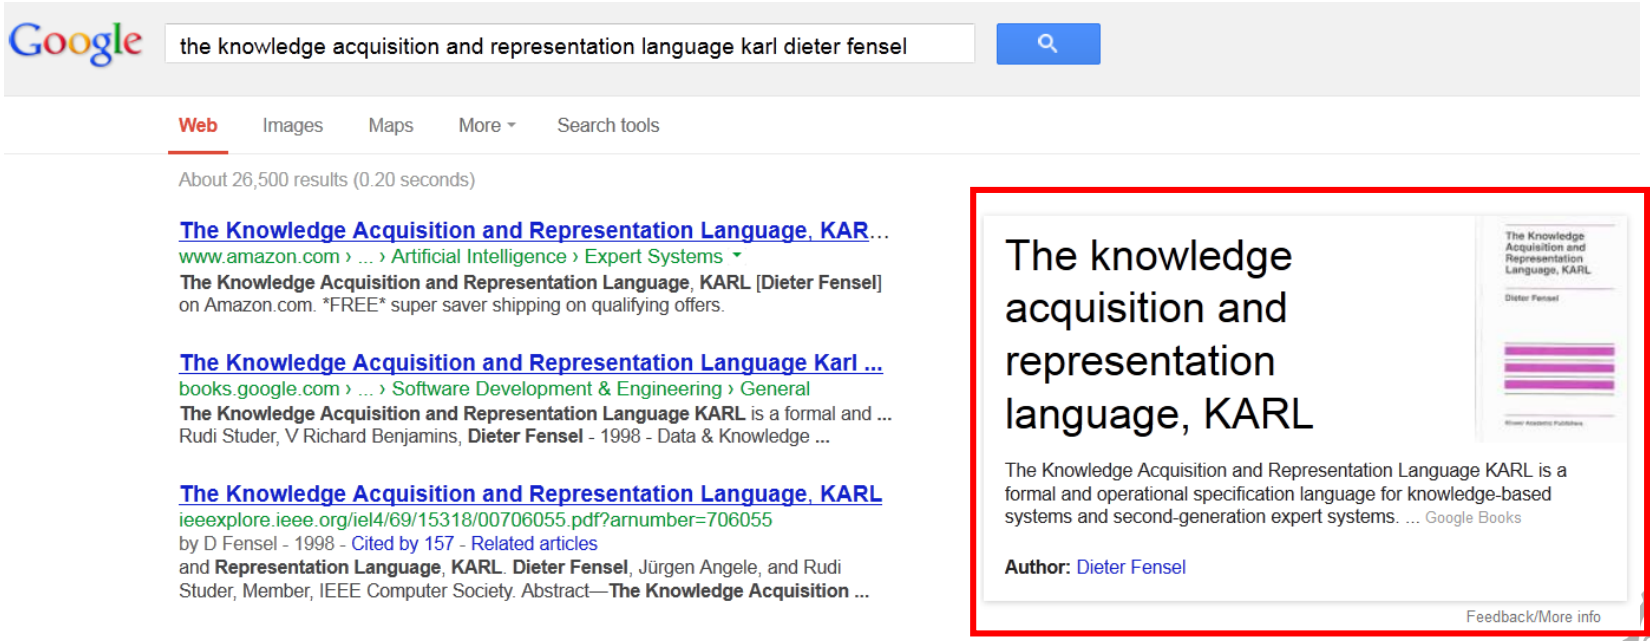

Sampling of GKG Entities

- Book
- BookSeries
- EducationalOrganization
- Event
- GovernmentOrganization
- LocalBusiness
- Movie
- MovieSeries
- MusicAlbum
- MusicGroup
- MusicRecording
- Organization
- Periodical
- Person
- Place
- SportsTeam
- TVEpisode
- TVSeries
- VideoGame
- VideoGameSeries
- WebSite

“Person” Schema

Property	Expected Type	Description
Properties from Person		
additionalName	Text	An additional name for a Person, can be used for a middle name.
address	PostalAddress or Text	Physical address of the item.
affiliation	Organization	An organization that this person is affiliated with. For example, a school/university, a club, or a team.
alumniOf	EducationalOrganization or Organization	An organization that the person is an alumni of. Inverse property: alumni.
award	Text	An award won by or for this item. Supersedes awards.
birthDate	Date	Date of birth.
birthPlace	Place	The place where the person was born.
brand	Brand or Organization	The brand(s) associated with a product or service, or the brand(s) maintained by an organization or business person.
children	Person	A child of the person.
colleague	Person or URL	A colleague of the person. Supersedes colleagues.
contactPoint	ContactPoint	A contact point for a person or organization. Supersedes contactPoints.
deathDate	Date	Date of death.
deathPlace	Place	The place where the person died.
duns	Text	The Dun & Bradstreet DUNS number for identifying an organization or business person.
email	Text	Email address.
familyName	Text	Family name. In the U.S., the last name of an Person. This can be used along with givenName instead of the name property.
faxNumber	Text	The fax number.
follows	Person	The most generic uni-directional social relation.

Data Sources Used for GKG

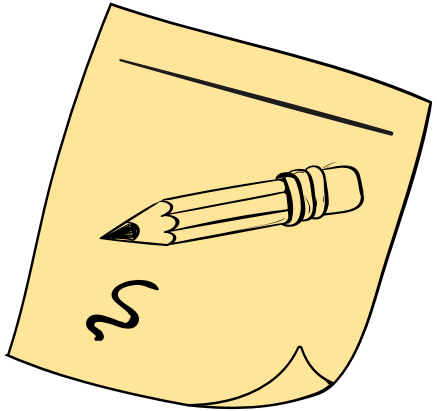

CIA World Factbook

- CIA World Factbook is a reference resource produced by the Central Intelligence Agency of the United States with almanac-style information about the countries of the world.

- GKG integrates information about geography, government, economy, etc. from CIA World Factbook.

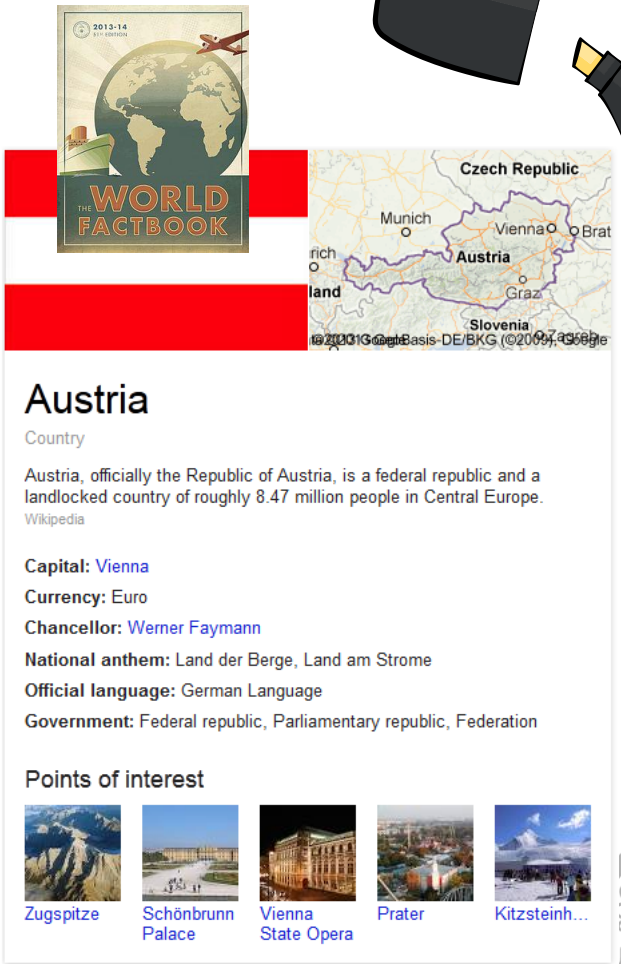

Freebase

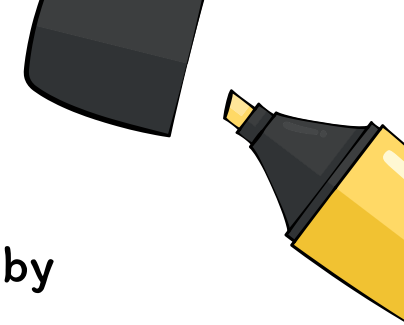

- Freebase is large collaborative knowledge base, developed by Metaweb and acquired by Google in 2010.

- GKG uses UIDs directly from the Freebase; detective work of Andreas Thalhammer showing how to get from GKG UIDs to Freebased UIDs using base64 and gzip

- Check the “Knowledge Graph links to Freebase” thread on w3c mailing list http://lists.w3.org/Archives/Public/semantic-web/2012Jun/0028.html

Wikipedia

- For most search results first sentences come from Wikipedia

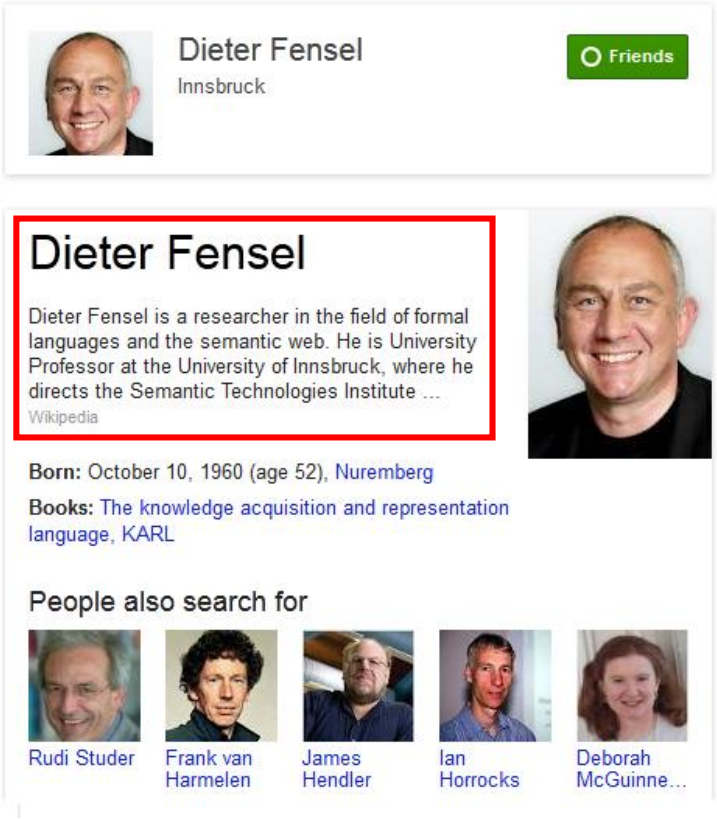

Other Sources

- GKG also considers the information Google retrieves from the volume of queries done by the users and the links those users have clicked on the results presented for those queries.

- GKG is also integrated with other Google products e.g. Google+

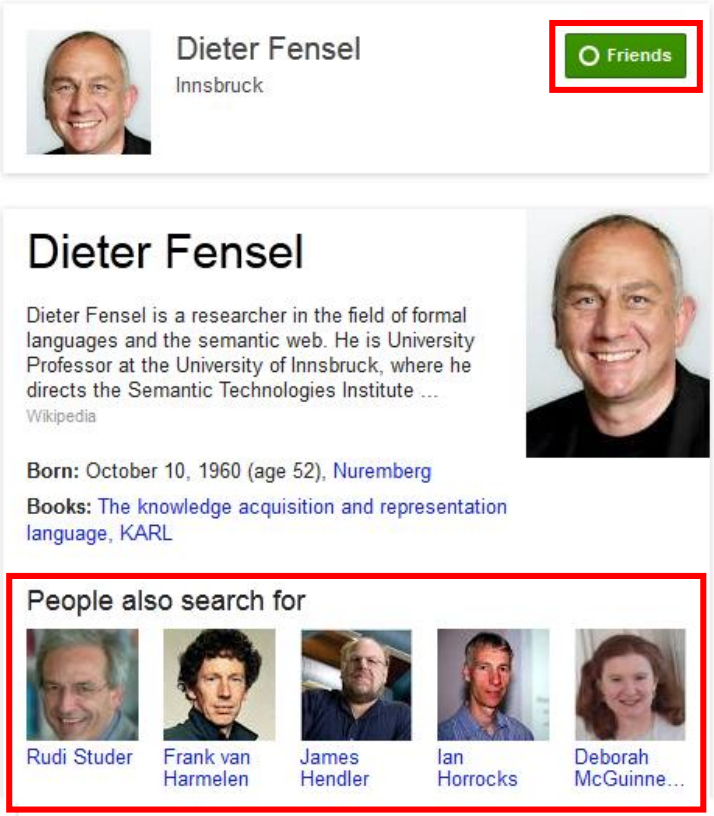

Web of Data

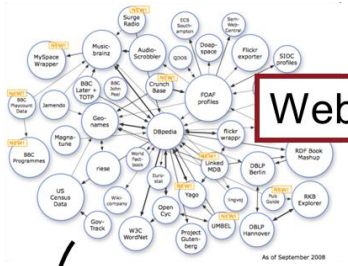

Picture from [4]

Web of Data

Semantic Web

Semantic Web

Semantic Annotations

Web

Hypermedia

Hypertext

?

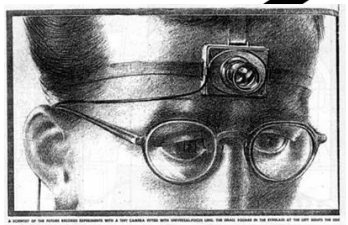

"As We May Think", 1945

Picture from http://www.theatlantic.com/doc/194507/bush

Web of Data

- Characteristics:
  - Links between arbitrary things (e.g., persons, locations, events, buildings)
  - Structure of data on Web pages is made explicit
  - Things described on Web pages are named and get URIs
  - Links between things are made explicit and are typed

- HTML defines how text should look when presented to humans. Semantic web markup defines how information should be organized to be interpretable by machines.

- Web of Data

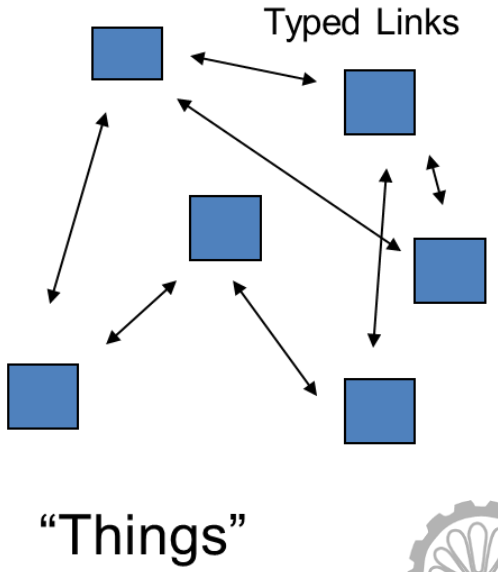

Web of Data

- A **closed** implementation of ***Web of Data*** principles
  - is not about documents, but objects such as people, places and things
  - objects are interlinked in the GKG
  - objects have structured information which is obtained from the web

- The Google Knowledge Graph is the basis for transforming Google' core search product from an *information engine* to a *knowledge engine (entity search engine)*

Building a KG Through ML

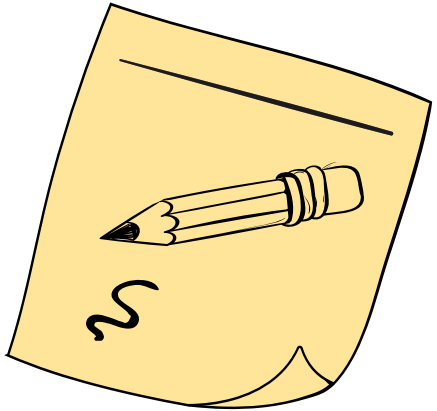

NELL: Never-Ending Language Learning

- Tom Mitchell et al. (CMU), 2010 to present.
- Learning to “read the web” 24 hours/day.
- Training data includes a collection of 1.2 billion web pages.
- Access to additional data through search engine APIs (100K calls/day).
- KB has 2.8 million instances over 1186 different categories.
- KB is freely available for download.
- You can help train NELL via Twitter.

Source: Never-Ending Learning (2018)

NELL Knowledge Fragment

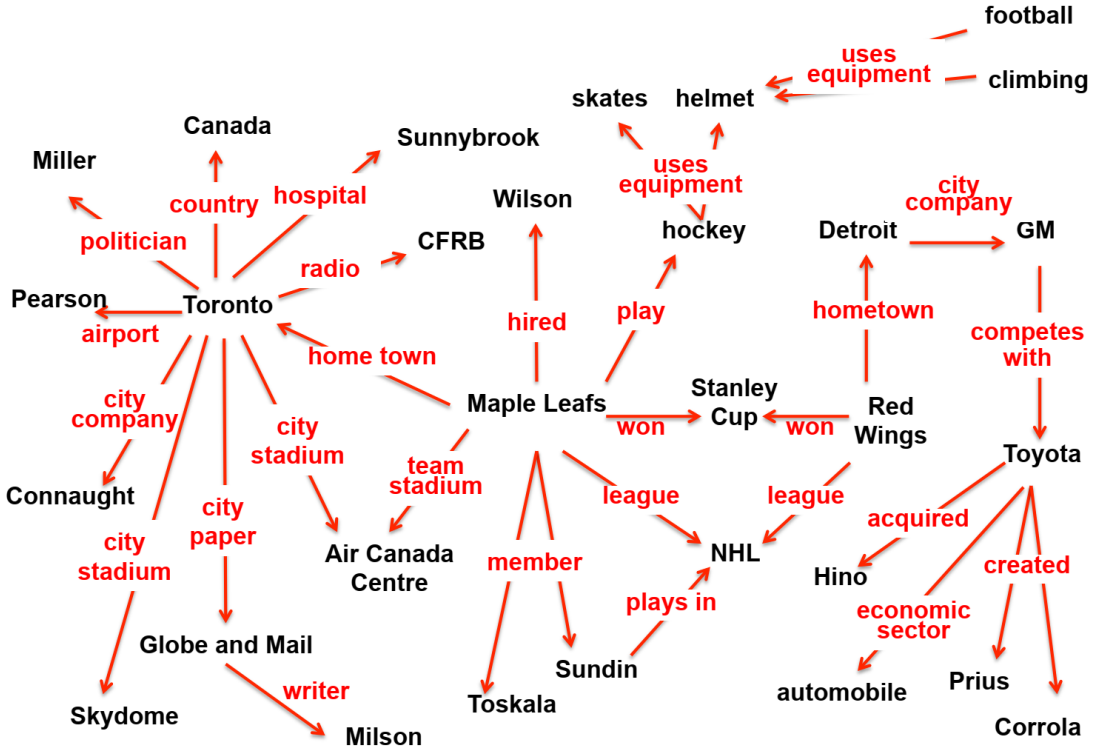

Motivation for NELL

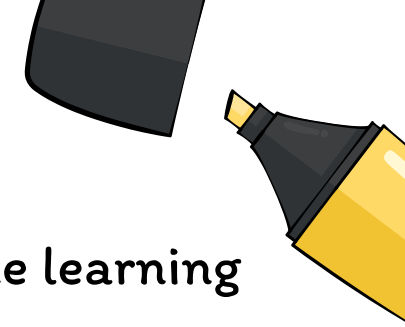

Thesis: “we will never truly understand human or machine learning until we can build computer programs that, like people,

- Learn many different types of knowledge or functions
- From years of diverse, mostly self-supervised experience
- In a staged curricular fashion, where previously learned knowledge enables learning further types of knowledge
- Where self-reflection and the ability to formulate new representations and new learning tasks enable the learning to avoid stagnation and performance plateaus.”

NELL Basic Idea

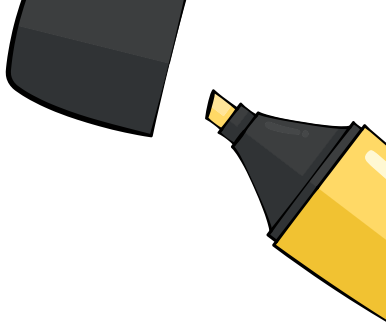

- NELL learns several things:
  - Categories
  - Triples: noun phrase 1 - relation - noun phrase 2
  - New relations
- Multiple inference algorithms propose triples and gather evidence for them.
  - Linguistic information
  - Word co-occurrence
  - Image labeling
  - Etc.
- Categories and triples supported by multiple sources of evidence grow in confidence.

NELL Architecture

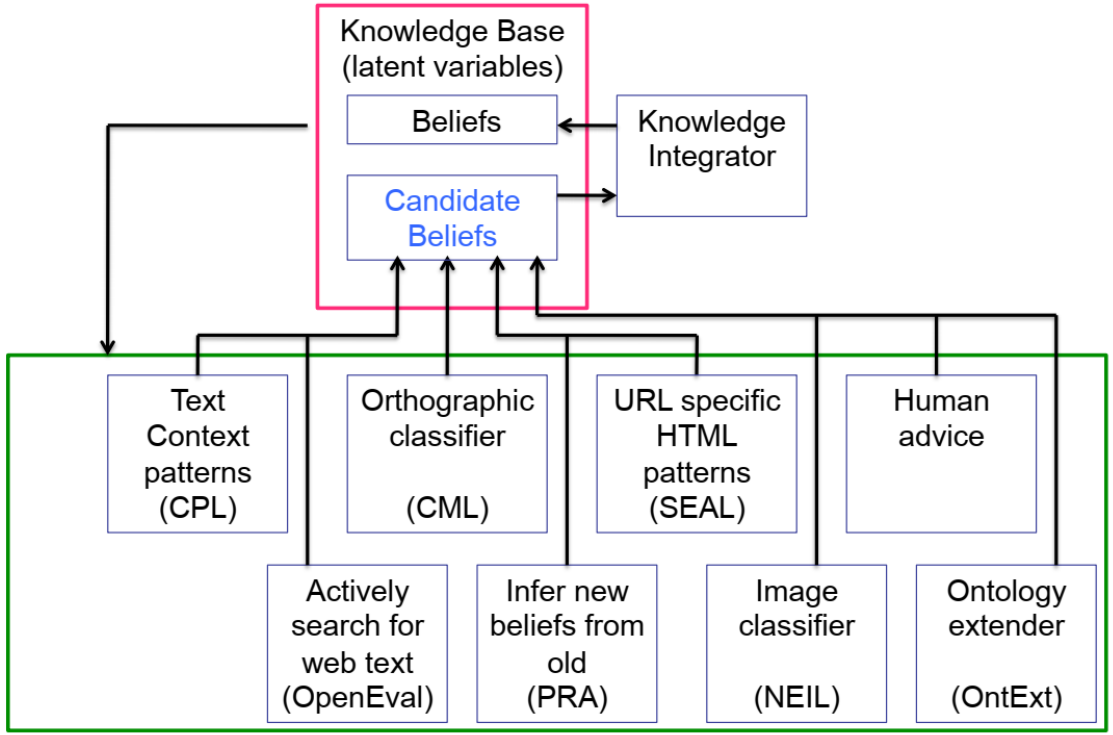

Never-Ending Learning

- Set of learning tasks $L = \{L_i\}$
- Task $L_i = < T_i, P_i, E_i >$
  - T_i is a task $< X_i, Y_i >$ specifying the domain of a function $f_i^* : X_i \rightarrow Y_i$
  - Pi is a performance metric P_i: $f \rightarrow \mathbb{R}$
  - E_i is an experience
- Coupling constraints $C = \{< \phi_k, V_k >\}$
  - ϕ_k specifies degree of satisfaction of the coupling constraint among tasks
  - V_k is a vector of indices over learning tasks specifying the arguments to ϕ_k
- $f_i^* = \min_{f \in F_i} P_i(f)$

✓ Goal: improve the quality of the task functions f_i as measured by the P_i.

✓ NELL faces over 4100 distinct learning tasks.

Category Classification Tasks

1. Character string features of the noun phrase: Coupled Morphological Classifier system (CMC)
2. Distribution of text contexts found around this noun phrase in the 1.2 billion page database: Coupled Pattern Learner system (CPL)
3. Distribution of text contexts found through active web search (OpenEval).
4. HTML structure of web pages that mention the noun phrase: Set Expander for Any Language system (SEAL)
5. Visual images associated with the noun phrase: Never Ending Image Learner (NEIL)
6. Learned vector embeddings (feature vectors) of the noun phrase: LE (Learned Embeddings)

Relation Classification

Does “Pittsburgh” + “US” satisfy the relation CityLocatedInCountry(x,y) ?

There are 461 relations in the ontology.

Four methods are used for relation classification:

1. Distribution of text contexts from CPL
2. Distribution of text context from OpenEval
3. HTML structure from SEAL
4. Learned vector embeddings from LE

Entity Resolution

- Functions to classify whether pairs of noun phrases are synonyms.
- Noun phrases are kept distinct from the entities to which they refer.
- Necessary to deal with polysemy.
  - “Coach” can be either a person or a vehicle.
- Two methods are used:
  - String similarity
  - Similarities in beliefs about the entities
- NELL learns for each category what are the good types of knowledge to take as evidence for synonymy.

Inference Rules Among Belief Triples

- Functions that propose new beliefs to be added to the KB.

- For each relation, the corresponding function is represented by a collection of restricted Horn Clause rules learned by the Path Ranking Algorithm (PRA)

Samples of Self-discovered Relations

- athleteWonAward
- animalEatsFood
- languageTaughtInCity
- clothingMadeFromPlant
- beverageServedWithFood
- fishServedWithFood
- athleteBeatAthlete
- athleteInjuredBodyPart
- arthropodFeedsOnInsect
- animalEatsVegetable
- plantRepresentsEmotion
- foodDecreasesRiskOfDisease
- clothingGoesWithClothing
- bacteriaCausesPhysCondition
- buildingFeatureMadeFromMaterial
- emotionAssociatedWithDisease
- foodCanCauseDisease
- agriculturalProductAttractsInsect
- arteryArisesFromArtery
- countryHasSportsFans
- bakedGoodServedWithBeverage
- beverageContainsProtein
- animalCanDevelopDisease
- beverageMadeFromBeverage

Growth of the KB Over Time

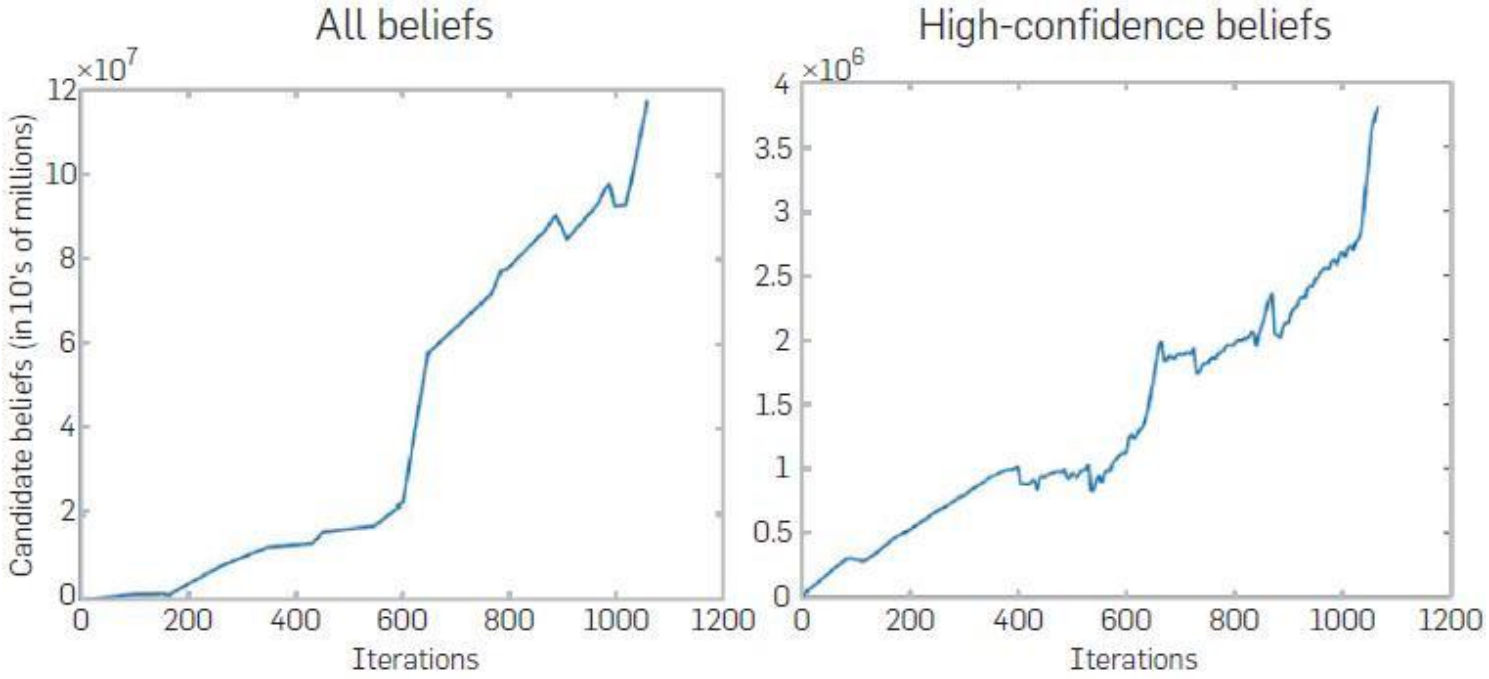

Human Correction of NELL

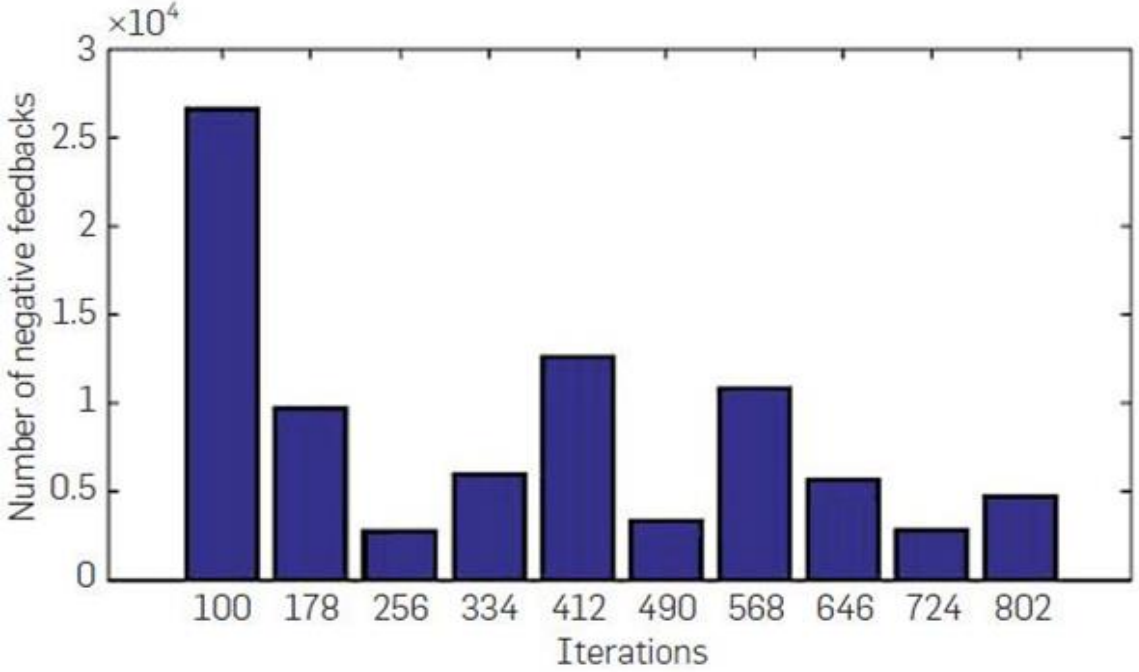

Average 2.4 negative feedback items per month per predicate.

Thanks for your attention

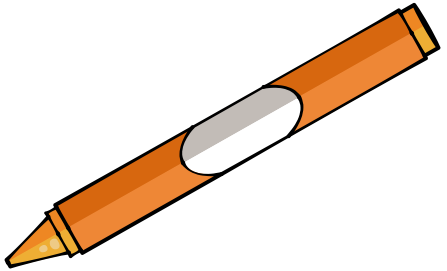

References and IP Notice

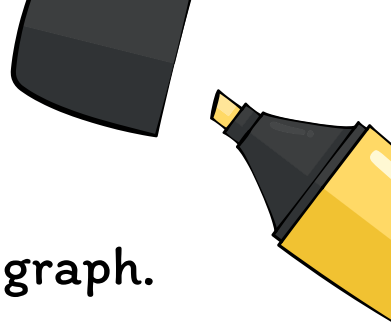

- Some slides from Dave Touretzky's slides on knowledge graph.
- The begging slides are mostly based on Ioan Toma’s slides.
- Some graphics from Slidesgo online template.